

recommend the one coauthored by James Gosling, the creator of the language [5]) describe the data formats and arithmetic operations supported by Java.

Most books on logic design [58, 116] have a section on encodings and arithmetic operations. Such books describe different ways of implementing arithmetic circuits. Overton's book on IEEE floating point [82] provides a detailed description of the format as well as the properties from the perspective of a numerical applications programmer.

## Homework Problems

### 2.55 ♦

Compile and run the sample code that uses `show_bytes` (file `show-bytes.c`) on different machines to which you have access. Determine the byte orderings used by these machines.

### 2.56 ♦

Try running the code for `show_bytes` for different sample values.

### 2.57 ♦

Write procedures `show_short`, `show_long`, and `show_double` that print the byte representations of C objects of types `short`, `long`, and `double`, respectively. Try these out on several machines.

### 2.58 ♦♦

Write a procedure `is_little_endian` that will return 1 when compiled and run on a little-endian machine, and will return 0 when compiled and run on a big-endian machine. This program should run on any machine, regardless of its word size.

### 2.59 ♦♦

Write a C expression that will yield a word consisting of the least significant byte of `x` and the remaining bytes of `y`. For operands `x = 0x89ABCDEF` and `y = 0x76543210`, this would give `0x765432EF`.

### 2.60 ♦♦

Suppose we number the bytes in a  $w$ -bit word from 0 (least significant) to  $w/8 - 1$  (most significant). Write code for the following C function, which will return an unsigned value in which byte `i` of argument `x` has been replaced by byte `b`:

```
unsigned replace_byte (unsigned x, int i, unsigned char b);
```

Here are some examples showing how the function should work:

```
replace_byte(0x12345678, 2, 0xAB) --> 0x12AB5678
replace_byte(0x12345678, 0, 0xAB) --> 0xAB345678
```

## Bit-Level Integer Coding Rules

In several of the following problems, we will artificially restrict what programming constructs you can use to help you gain a better understanding of the bit-level,

logic, and arithmetic operations of C. In answering these problems, your code must follow these rules:

- Assumptions
  - Integers are represented in two's-complement form.
  - Right shifts of signed data are performed arithmetically.
  - Data type `int` is  $w$  bits long. For some of the problems, you will be given a specific value for  $w$ , but otherwise your code should work as long as  $w$  is a multiple of 8. You can use the expression `sizeof(int)<<3` to compute  $w$ .
- Forbidden
  - Conditionals (`if` or `?:`), loops, switch statements, function calls, and macro invocations.
  - Division, modulus, and multiplication.
  - Relative comparison operators (`<`, `>`, `<=`, and `>=`).
- Allowed operations
  - All bit-level and logic operations.
  - Left and right shifts, but only with shift amounts between 0 and  $w - 1$ .
  - Addition and subtraction.
  - Equality (`==`) and inequality (`!=`) tests. (Some of the problems do not allow these.)
  - Integer constants `INT_MIN` and `INT_MAX`.
  - Casting between data types `int` and `unsigned`, either explicitly or implicitly.

Even with these rules, you should try to make your code readable by choosing descriptive variable names and using comments to describe the logic behind your solutions. As an example, the following code extracts the most significant byte from integer argument `x`:

```
/* Get most significant byte from x */
int get_msb(int x) {
    /* Shift by w-8 */
    int shift_val = (sizeof(int)-1)<<3;
    /* Arithmetic shift */
    int xright = x >> shift_val;
    /* Zero all but LSB */
    return xright & 0xFF;
}
```

## 2.61 ♦♦

Write C expressions that evaluate to 1 when the following conditions are true and to 0 when they are false. Assume `x` is of type `int`.

- A. Any bit of `x` equals 1.
- B. Any bit of `x` equals 0.

- C. Any bit in the least significant byte of  $x$  equals 1.
- D. Any bit in the most significant byte of  $x$  equals 0.

Your code should follow the bit-level integer coding rules (page 164), with the additional restriction that you may not use equality (`==`) or inequality (`!=`) tests.

### 2.62 ♦♦♦

Write a function `int_shifts_are_arithmetic()` that yields 1 when run on a machine that uses arithmetic right shifts for data type `int` and yields 0 otherwise. Your code should work on a machine with any word size. Test your code on several machines.

### 2.63 ♦♦♦

Fill in code for the following C functions. Function `srl` performs a logical right shift using an arithmetic right shift (given by value `xsra`), followed by other operations not including right shifts or division. Function `sra` performs an arithmetic right shift using a logical right shift (given by value `xsrl`), followed by other operations not including right shifts or division. You may use the computation `8*sizeof(int)` to determine  $w$ , the number of bits in data type `int`. The shift amount  $k$  can range from 0 to  $w - 1$ .

```
unsigned srl(unsigned x, int k) {
    /* Perform shift arithmetically */
    unsigned xsra = (int) x >> k;
    .
    .
    .
    .
}

int sra(int x, int k) {
    /* Perform shift logically */
    int xsrl = (unsigned) x >> k;
    .
    .
    .
    .
}
```

### 2.64 ♦

Write code to implement the following function:

```
/* Return 1 when any odd bit of x equals 1; 0 otherwise.
   Assume w=32 */
int any_odd_one(unsigned x);
```

Your function should follow the bit-level integer coding rules (page 164), except that you may assume that data type `int` has  $w = 32$  bits.

## 2.65 ♦♦♦♦

Write code to implement the following function:

```
/* Return 1 when x contains an odd number of 1s; 0 otherwise.
   Assume w=32 */
int odd_ones(unsigned x);
```

Your function should follow the bit-level integer coding rules (page 164), except that you may assume that data type `int` has  $w = 32$  bits.

Your code should contain a total of at most 12 arithmetic, bitwise, and logical operations.

## 2.66 ♦♦♦

Write code to implement the following function:

```
/*
 * Generate mask indicating leftmost 1 in x. Assume w=32.
 * For example, 0xFF00 -> 0x8000, and 0x6600 --> 0x4000.
 * If x = 0, then return 0.
 */
int leftmost_one(unsigned x);
```

Your function should follow the bit-level integer coding rules (page 164), except that you may assume that data type `int` has  $w = 32$  bits.

Your code should contain a total of at most 15 arithmetic, bitwise, and logical operations.

*Hint:* First transform  $x$  into a bit vector of the form  $[0 \dots 011 \dots 1]$ .

## 2.67 ♦♦

You are given the task of writing a procedure `int_size_is_32()` that yields 1 when run on a machine for which an `int` is 32 bits, and yields 0 otherwise. You are not allowed to use the `sizeof` operator. Here is a first attempt:

```
1  /* The following code does not run properly on some machines */
2  int bad_int_size_is_32() {
3      /* Set most significant bit (msb) of 32-bit machine */
4      int set_msb = 1 << 31;
5      /* Shift past msb of 32-bit word */
6      int beyond_msb = 1 << 32;
7
8      /* set_msb is nonzero when word size >= 32
9       beyond_msb is zero when word size <= 32 */
10     return set_msb && !beyond_msb;
11 }
```

When compiled and run on a 32-bit SUN SPARC, however, this procedure returns 0. The following compiler message gives us an indication of the problem:

```
warning: left shift count >= width of type
```

- A. In what way does our code fail to comply with the C standard?
- B. Modify the code to run properly on any machine for which data type `int` is at least 32 bits.
- C. Modify the code to run properly on any machine for which data type `int` is at least 16 bits.

**2.68 ♦♦**

Write code for a function with the following prototype:

```
/*
 * Mask with least significant n bits set to 1
 * Examples: n = 6 --> 0x3F, n = 17 --> 0x1FFFF
 * Assume 1 <= n <= w
 */
int lower_one_mask(int n);
```

Your function should follow the bit-level integer coding rules (page 164). Be careful of the case  $n = w$ .

**2.69 ♦♦♦**

Write code for a function with the following prototype:

```
/*
 * Do rotating left shift. Assume 0 <= n < w
 * Examples when x = 0x12345678 and w = 32:
 *   n=4 -> 0x23456781, n=20 -> 0x67812345
 */
unsigned rotate_left(unsigned x, int n);
```

Your function should follow the bit-level integer coding rules (page 164). Be careful of the case  $n = 0$ .

**2.70 ♦♦**

Write code for the function with the following prototype:

```
/*
 * Return 1 when x can be represented as an n-bit, 2's-complement
 * number; 0 otherwise
 * Assume 1 <= n <= w
 */
int fits_bits(int x, int n);
```

Your function should follow the bit-level integer coding rules (page 164).

**2.71 ♦**

You just started working for a company that is implementing a set of procedures to operate on a data structure where 4 signed bytes are packed into a 32-bit unsigned. Bytes within the word are numbered from 0 (least significant) to 3

(most significant). You have been assigned the task of implementing a function for a machine using two's-complement arithmetic and arithmetic right shifts with the following prototype:

```
/* Declaration of data type where 4 bytes are packed
   into an unsigned */
typedef unsigned packed_t;

/* Extract byte from word. Return as signed integer */
int xbyte(packed_t word, int bytenum);
```

That is, the function will extract the designated byte and sign extend it to be a 32-bit int.

Your predecessor (who was fired for incompetence) wrote the following code:

```
/* Failed attempt at xbyte */
int xbyte(packed_t word, int bytenum)
{
    return (word >> (bytenum << 3)) & 0xFF;
}
```

- A. What is wrong with this code?
- B. Give a correct implementation of the function that uses only left and right shifts, along with one subtraction.

## 2.72 ♦♦

You are given the task of writing a function that will copy an integer `val` into a buffer `buf`, but it should do so only if enough space is available in the buffer.

Here is the code you write:

```
/* Copy integer into buffer if space is available */
/* WARNING: The following code is buggy */
void copy_int(int val, void *buf, int maxbytes) {
    if (maxbytes - sizeof(val) >= 0)
        memcpy(buf, (void *) &val, sizeof(val));
}
```

This code makes use of the library function `memcpy`. Although its use is a bit artificial here, where we simply want to copy an `int`, it illustrates an approach commonly used to copy larger data structures.

You carefully test the code and discover that it *always* copies the value to the buffer, even when `maxbytes` is too small.

- A. Explain why the conditional test in the code always succeeds. *Hint:* The `sizeof` operator returns a value of type `size_t`.
- B. Show how you can rewrite the conditional test to make it work properly.

**2.73 ♦♦**

Write code for a function with the following prototype:

```
/* Addition that saturates to TMin or TMax */
int saturating_add(int x, int y);
```

Instead of overflowing the way normal two's-complement addition does, saturating addition returns *TMax* when there would be positive overflow, and *TMin* when there would be negative overflow. Saturating arithmetic is commonly used in programs that perform digital signal processing.

Your function should follow the bit-level integer coding rules (page 164).

**2.74 ♦♦**

Write a function with the following prototype:

```
/* Determine whether arguments can be subtracted without overflow */
int tsub_ok(int x, int y);
```

This function should return 1 if the computation  $x - y$  does not overflow.

**2.75 ♦♦♦**

Suppose we want to compute the complete  $2w$ -bit representation of  $x \cdot y$ , where both  $x$  and  $y$  are unsigned, on a machine for which data type unsigned is  $w$  bits. The low-order  $w$  bits of the product can be computed with the expression  $x * y$ , so we only require a procedure with prototype

```
unsigned unsigned_high_prod(unsigned x, unsigned y);
```

that computes the high-order  $w$  bits of  $x \cdot y$  for unsigned variables.

We have access to a library function with prototype

```
int signed_high_prod(int x, int y);
```

that computes the high-order  $w$  bits of  $x \cdot y$  for the case where  $x$  and  $y$  are in two's-complement form. Write code calling this procedure to implement the function for unsigned arguments. Justify the correctness of your solution.

*Hint:* Look at the relationship between the signed product  $x \cdot y$  and the unsigned product  $x' \cdot y'$  in the derivation of Equation 2.18.

**2.76 ♦**

The library function `calloc` has the following declaration:

```
void *calloc(size_t nmemb, size_t size);
```

According to the library documentation, “The `calloc` function allocates memory for an array of `nmemb` elements of size bytes each. The memory is set to zero. If `nmemb` or `size` is zero, then `calloc` returns `NULL`.”

Write an implementation of `calloc` that performs the allocation by a call to `malloc` and sets the memory to zero via `memset`. Your code should not have any vulnerabilities due to arithmetic overflow, and it should work correctly regardless of the number of bits used to represent data of type `size_t`.

As a reference, functions `malloc` and `memset` have the following declarations:

```
void *malloc(size_t size);
void *memset(void *s, int c, size_t n);
```

**2.77 ♦♦**

Suppose we are given the task of generating code to multiply integer variable  $x$  by various different constant factors  $K$ . To be efficient, we want to use only the operations  $+$ ,  $-$ , and  $\ll$ . For the following values of  $K$ , write C expressions to perform the multiplication using at most three operations per expression.

- A.  $K = 17$
- B.  $K = -7$
- C.  $K = 60$
- D.  $K = -112$

**2.78 ♦♦**

Write code for a function with the following prototype:

```
/* Divide by power of 2. Assume 0 <= k < w-1 */
int divide_power2(int x, int k);
```

The function should compute  $x/2^k$  with correct rounding, and it should follow the bit-level integer coding rules (page 164).

**2.79 ♦♦**

Write code for a function `mul3div4` that, for integer argument  $x$ , computes  $3 * x/4$  but follows the bit-level integer coding rules (page 164). Your code should replicate the fact that the computation  $3*x$  can cause overflow.

**2.80 ♦♦♦**

Write code for a function `threefourths` that, for integer argument  $x$ , computes the value of  $\frac{3}{4}x$ , rounded toward zero. It should not overflow. Your function should follow the bit-level integer coding rules (page 164).

**2.81 ♦♦**

Write C expressions to generate the bit patterns that follow, where  $a^k$  represents  $k$  repetitions of symbol  $a$ . Assume a  $w$ -bit data type. Your code may contain references to parameters  $j$  and  $k$ , representing the values of  $j$  and  $k$ , but not a parameter representing  $w$ .

- A.  $1^{w-k}0^k$
- B.  $0^{w-k-j}1^k0^j$

**2.82 ♦**

We are running programs where values of type `int` are 32 bits. They are represented in two's complement, and they are right shifted arithmetically. Values of type `unsigned` are also 32 bits.



We generate arbitrary values  $x$  and  $y$ , and convert them to unsigned values as follows:

```
/* Create some arbitrary values */
int x = random();
int y = random();
/* Convert to unsigned */
unsigned ux = (unsigned) x;
unsigned uy = (unsigned) y;
```

For each of the following C expressions, you are to indicate whether or not the expression *always* yields 1. If it always yields 1, describe the underlying mathematical principles. Otherwise, give an example of arguments that make it yield 0.

- A.  $(x < y) == (-x > -y)$
- B.  $((x + y) < 4) + y - x == 17 * y + 15 * x$
- C.  $\sim x + \sim y + 1 == \sim(x + y)$
- D.  $(ux - uy) == -(\text{unsigned})(y - x)$
- E.  $((x >> 2) << 2) <= x$

### 2.83 ♦♦

Consider numbers having a binary representation consisting of an infinite string of the form  $0.yyy\ldots$ , where  $y$  is a  $k$ -bit sequence. For example, the binary representation of  $\frac{1}{3}$  is  $0.01010101\ldots$  ( $y = 01$ ), while the representation of  $\frac{1}{5}$  is  $0.001100110011\ldots$  ( $y = 0011$ ).

- A. Let  $Y = B2U_k(y)$ , that is, the number having binary representation  $y$ . Give a formula in terms of  $Y$  and  $k$  for the value represented by the infinite string. *Hint:* Consider the effect of shifting the binary point  $k$  positions to the right.
- B. What is the numeric value of the string for the following values of  $y$ ?
  - (a) 101
  - (b) 0110
  - (c) 010011

### 2.84 ♦

Fill in the return value for the following procedure, which tests whether its first argument is less than or equal to its second. Assume the function `f2u` returns an unsigned 32-bit number having the same bit representation as its floating-point argument. You can assume that neither argument is *NaN*. The two flavors of zero,  $+0$  and  $-0$ , are considered equal.

```
int float_le(float x, float y) {
    unsigned ux = f2u(x);
    unsigned uy = f2u(y);
```

```

/* Get the sign bits */
unsigned sx = ux >> 31;
unsigned sy = uy >> 31;

/* Give an expression using only ux, uy, sx, and sy */
return      ;
}

```

**2.85** ◆

Given a floating-point format with a  $k$ -bit exponent and an  $n$ -bit fraction, write formulas for the exponent  $E$ , the significand  $M$ , the fraction  $f$ , and the value  $V$  for the quantities that follow. In addition, describe the bit representation.

- A. The number 7.0
- B. The largest odd integer that can be represented exactly
- C. The reciprocal of the smallest positive normalized value

**2.86** ◆

Intel-compatible processors also support an “extended-precision” floating-point format with an 80-bit word divided into a sign bit,  $k = 15$  exponent bits, a single *integer* bit, and  $n = 63$  fraction bits. The integer bit is an explicit copy of the implied bit in the IEEE floating-point representation. That is, it equals 1 for normalized values and 0 for denormalized values. Fill in the following table giving the approximate values of some “interesting” numbers in this format:

Description	Extended precision	
	Value	Decimal
Smallest positive denormalized	_____	_____
Smallest positive normalized	_____	_____
Largest normalized	_____	_____

This format can be used in C programs compiled for Intel-compatible machines by declaring the data to be of type `long double`. However, it forces the compiler to generate code based on the legacy 8087 floating-point instructions. The resulting program will most likely run much slower than would be the case for data type `float` or `double`.

**2.87** ◆

The 2008 version of the IEEE floating-point standard, named IEEE 754-2008, includes a 16-bit “half-precision” floating-point format. It was originally devised by computer graphics companies for storing data in which a higher dynamic range is required than can be achieved with 16-bit integers. This format has 1 sign bit, 5 exponent bits ( $k = 5$ ), and 10 fraction bits ( $n = 10$ ). The exponent bias is  $2^{5-1} - 1 = 15$ .

Fill in the table that follows for each of the numbers given, with the following instructions for each column:

Hex: The four hexadecimal digits describing the encoded form.

$M$ : The value of the significand. This should be a number of the form  $x$  or  $\frac{x}{y}$ , where  $x$  is an integer and  $y$  is an integral power of 2. Examples include 0,  $\frac{67}{64}$ , and  $\frac{1}{256}$ .

$E$ : The integer value of the exponent.

$V$ : The numeric value represented. Use the notation  $x$  or  $x \times 2^z$ , where  $x$  and  $z$  are integers.

$D$ : The (possibly approximate) numerical value, as is printed using the `%f` formatting specification of `printf`.

As an example, to represent the number  $\frac{7}{8}$ , we would have  $s = 0$ ,  $M = \frac{7}{4}$ , and  $E = -1$ . Our number would therefore have an exponent field of  $01110_2$  (decimal value  $15 - 1 = 14$ ) and a significand field of  $1100000000_2$ , giving a hex representation 3B00. The numerical value is 0.875.

You need not fill in entries marked —.

Description	Hex	$M$	$E$	$V$	$D$
−0	_____	_____	_____	−0	−0.0
Smallest value > 2	_____	_____	_____	_____	_____
512	_____	_____	_____	_____512	512.0
Largest denormalized	_____	_____	_____	_____	_____
−∞	_____	—	—	−∞	−∞
Number with hex representation 3BB0	3BB0	_____	_____	_____	_____

## 2.88 ♦♦

Consider the following two 9-bit floating-point representations based on the IEEE floating-point format.

### 1. Format A

- There is 1 sign bit.
- There are  $k = 5$  exponent bits. The exponent bias is 15.
- There are  $n = 3$  fraction bits.

### 2. Format B

- There is 1 sign bit.
- There are  $k = 4$  exponent bits. The exponent bias is 7.
- There are  $n = 4$  fraction bits.

In the following table, you are given some bit patterns in format A, and your task is to convert them to the closest value in format B. If rounding is necessary you should *round toward*  $+\infty$ . In addition, give the values of numbers given by the format A and format B bit patterns. Give these as whole numbers (e.g., 17) or as fractions (e.g.,  $17/64$  or  $17/2^6$ ).

Format A		Format B	
Bits	Value	Bits	Value
1 01111 001	$-\frac{9}{8}$	1 0111 0010	$-\frac{9}{8}$
0 10110 011	_____	_____	_____
1 00111 010	_____	_____	_____
0 00000 111	_____	_____	_____
1 11100 000	_____	_____	_____
0 10111 100	_____	_____	_____

**2.89** ◆

We are running programs on a machine where values of type `int` have a 32-bit two's-complement representation. Values of type `float` use the 32-bit IEEE format, and values of type `double` use the 64-bit IEEE format.

We generate arbitrary integer values `x`, `y`, and `z`, and convert them to values of type `double` as follows:

```
/* Create some arbitrary values */
int x = random();
int y = random();
int z = random();
/* Convert to double */
double dx = (double) x;
double dy = (double) y;
double dz = (double) z;
```

For each of the following C expressions, you are to indicate whether or not the expression *always* yields 1. If it always yields 1, describe the underlying mathematical principles. Otherwise, give an example of arguments that make it yield 0. Note that you cannot use an IA32 machine running gcc to test your answers, since it would use the 80-bit extended-precision representation for both `float` and `double`.

- A. `(float) x == (float) dx`
- B. `dx - dy == (double) (x-y)`
- C. `(dx + dy) + dz == dx + (dy + dz)`
- D. `(dx * dy) * dz == dx * (dy * dz)`
- E. `dx / dx == dz / dz`

**2.90** ◆

You have been assigned the task of writing a C function to compute a floating-point representation of  $2^x$ . You decide that the best way to do this is to directly construct the IEEE single-precision representation of the result. When  $x$  is too small, your routine will return 0.0. When  $x$  is too large, it will return  $+\infty$ . Fill in the blank portions of the code that follows to compute the correct result. Assume the

function `u2f` returns a floating-point value having an identical bit representation as its unsigned argument.

```
float fpwr2(int x)
{
    /* Result exponent and fraction */
    unsigned exp, frac;
    unsigned u;

    if (x < _____) {
        /* Too small. Return 0.0 */
        exp = _____;
        frac = _____;
    } else if (x < _____) {
        /* Denormalized result */
        exp = _____;
        frac = _____;
    } else if (x < _____) {
        /* Normalized result. */
        exp = _____;
        frac = _____;
    } else {
        /* Too big. Return +oo */
        exp = _____;
        frac = _____;
    }

    /* Pack exp and frac into 32 bits */
    u = exp << 23 | frac;
    /* Return as float */
    return u2f(u);
}
```

### 2.91 ♦

Around 250 B.C., the Greek mathematician Archimedes proved that  $\frac{223}{71} < \pi < \frac{22}{7}$ . Had he had access to a computer and the standard library `<math.h>`, he would have been able to determine that the single-precision floating-point approximation of  $\pi$  has the hexadecimal representation `0x40490FDB`. Of course, all of these are just approximations, since  $\pi$  is not rational.

- A. What is the fractional binary number denoted by this floating-point value?
- B. What is the fractional binary representation of  $\frac{22}{7}$ ? *Hint:* See Problem 2.83.
- C. At what bit position (relative to the binary point) do these two approximations to  $\pi$  diverge?

### Bit-Level Floating-Point Coding Rules

In the following problems, you will write code to implement floating-point functions, operating directly on bit-level representations of floating-point numbers. Your code should exactly replicate the conventions for IEEE floating-point operations, including using round-to-even mode when rounding is required.

To this end, we define data type `float_bits` to be equivalent to unsigned:

```
/* Access bit-level representation floating-point number */
typedef unsigned float_bits;
```

Rather than using data type `float` in your code, you will use `float_bits`. You may use both `int` and unsigned data types, including unsigned and integer constants and operations. You may not use any unions, structs, or arrays. Most significantly, you may not use any floating-point data types, operations, or constants. Instead, your code should perform the bit manipulations that implement the specified floating-point operations.

The following function illustrates the use of these coding rules. For argument  $f$ , it returns  $\pm 0$  if  $f$  is denormalized (preserving the sign of  $f$ ), and returns  $f$  otherwise.

```
/* If f is denorm, return 0. Otherwise, return f */
float_bits float_denorm_zero(float_bits f) {
    /* Decompose bit representation into parts */
    unsigned sign = f >> 31;
    unsigned exp = f >> 23 & 0xFF;
    unsigned frac = f & 0x7FFFFFFF;
    if (exp == 0) {
        /* Denormalized. Set fraction to 0 */
        frac = 0;
    }
    /* Reassemble bits */
    return (sign << 31) | (exp << 23) | frac;
}
```

### 2.92 ♦♦

Following the bit-level floating-point coding rules, implement the function with the following prototype:

```
/* Compute -f. If f is NaN, then return f. */
float_bits float_negate(float_bits f);
```

For floating-point number  $f$ , this function computes  $-f$ . If  $f$  is *NaN*, your function should simply return  $f$ .

Test your function by evaluating it for all  $2^{32}$  values of argument  $f$  and comparing the result to what would be obtained using your machine's floating-point operations.

**2.93 ♦♦**

Following the bit-level floating-point coding rules, implement the function with the following prototype:

```
/* Compute |f|. If f is NaN, then return f. */
float_bits float_absval(float_bits f);
```

For floating-point number  $f$ , this function computes  $|f|$ . If  $f$  is *NaN*, your function should simply return  $f$ .

Test your function by evaluating it for all  $2^{32}$  values of argument  $f$  and comparing the result to what would be obtained using your machine's floating-point operations.

**2.94 ♦♦♦**

Following the bit-level floating-point coding rules, implement the function with the following prototype:

```
/* Compute 2*f. If f is NaN, then return f. */
float_bits float_twice(float_bits f);
```

For floating-point number  $f$ , this function computes  $2.0 \cdot f$ . If  $f$  is *NaN*, your function should simply return  $f$ .

Test your function by evaluating it for all  $2^{32}$  values of argument  $f$  and comparing the result to what would be obtained using your machine's floating-point operations.

**2.95 ♦♦♦**

Following the bit-level floating-point coding rules, implement the function with the following prototype:

```
/* Compute 0.5*f. If f is NaN, then return f. */
float_bits float_half(float_bits f);
```

For floating-point number  $f$ , this function computes  $0.5 \cdot f$ . If  $f$  is *NaN*, your function should simply return  $f$ .

Test your function by evaluating it for all  $2^{32}$  values of argument  $f$  and comparing the result to what would be obtained using your machine's floating-point operations.

**2.96 ♦♦♦♦**

Following the bit-level floating-point coding rules, implement the function with the following prototype:

```
/*
 * Compute (int) f.
 * If conversion causes overflow or f is NaN, return 0x80000000
 */
int float_f2i(float_bits f);
```

For floating-point number  $f$ , this function computes  $(\text{int}) f$ . Your function should round toward zero. If  $f$  cannot be represented as an integer (e.g., it is out of range, or it is *NaN*), then the function should return `0x80000000`.

Test your function by evaluating it for all  $2^{32}$  values of argument  $f$  and comparing the result to what would be obtained using your machine's floating-point operations.

### 2.97 ♦♦♦♦

Following the bit-level floating-point coding rules, implement the function with the following prototype:

```
/* Compute (float) i */
float_bits float_i2f(int i);
```

For argument  $i$ , this function computes the bit-level representation of  $(\text{float}) i$ .

Test your function by evaluating it for all  $2^{32}$  values of argument  $f$  and comparing the result to what would be obtained using your machine's floating-point operations.

## Solutions to Practice Problems

### Solution to Problem 2.1 (page 73)

Understanding the relation between hexadecimal and binary formats will be important once we start looking at machine-level programs. The method for doing these conversions is in the text, but it takes a little practice to become familiar.

A. `0x25B9D2` to binary:

Hexadecimal	2	5	B	9	D	2
Binary	0010	0101	1101	1001	1101	0010

B. Binary `1100 1001 0111 1011` to hexadecimal:

Binary	1100	1001	0111	1011
Hexadecimal	C	9	7	B

C. `0xA8B3D` to binary:

Hexadecimal	A	8	B	3	D
Binary	1010	1000	1011	0011	1101

D. Binary `11 0010 0010 1101 1001 0110` to hexadecimal:

Binary	11	0010	0010	1101	1001	0110
Hexadecimal	3	2	2	D	9	6

### Solution to Problem 2.2 (page 73)

This problem gives you a chance to think about powers of 2 and their hexadecimal representations.



$n$	$2^n$ (decimal)	$2^n$ (hexadecimal)
5	32	0x20
23	8,388,608	0x800000
15	32,768	0x8000
13	8,192	0x2000
12	4,096	0x1000
6	64	0x40
8	256	0x100

**Solution to Problem 2.3 (page 74)**

This problem gives you a chance to try out conversions between hexadecimal and decimal representations for some smaller numbers. For larger ones, it becomes much more convenient and reliable to use a calculator or conversion program.

Decimal	Binary	Hexadecimal
0	0000 0000	0x00
$158 = 16 \cdot 9 + 14$	1001 1110	0x9E
$76 = 16 \cdot 4 + 12$	0100 1100	0x4C
$145 = 16 \cdot 9 + 1$	1001 0001	0x91
$16 \cdot 10 + 14 = 174$	1010 1110	0xAE
$16 \cdot 3 + 12 = 60$	0011 1100	0x3C
$16 \cdot 15 + 1 = 241$	1111 0001	0xF1
$16 \cdot 7 + 5 = 117$	0111 0101	0x75
$16 \cdot 11 + 13 = 189$	1011 1101	0xBD
$16 \cdot 15 + 5 = 245$	1111 0101	0xF5

**Solution to Problem 2.4 (page 75)**

When you begin debugging machine-level programs, you will find many cases where some simple hexadecimal arithmetic would be useful. You can always convert numbers to decimal, perform the arithmetic, and convert them back, but being able to work directly in hexadecimal is more efficient and informative.

- A.  $0x605C + 0x5 = 0x6061$ . Adding 5 to hex C gives 1 with a carry of 1.
- B.  $0x605C - 0x20 = 0x603C$ . Subtracting 2 from 5 in the second digit position requires no borrow from the third. This gives 3.
- C.  $0x605C + 32 = 0x607C$ . Decimal 32 ( $2^5$ ) equals hexadecimal 0x20.
- D.  $0x60FA - 0x605C = 0x9E$ . To subtract hex C (decimal 12) from hex A (decimal 10), we borrow 16 from the second digit, giving hex F (decimal 15). In the second digit, we now subtract 5 from hex E (decimal 14), giving decimal 9.

**Solution to Problem 2.5 (page 84)**

This problem tests your understanding of the byte representation of data and the two different byte orderings.

- A. Little endian: 78                      Big endian: 12
- B. Little endian: 78 56                  Big endian: 12 34

C. Little endian: 78 56 34      Big endian: 12 34 56

Recall that `show_bytes` enumerates a series of bytes starting from the one with lowest address and working toward the one with highest address. On a little-endian machine, it will list the bytes from least significant to most. On a big-endian machine, it will list bytes from the most significant byte to the least.

#### Solution to Problem 2.6 (page 85)

This problem is another chance to practice hexadecimal to binary conversion. It also gets you thinking about integer and floating-point representations. We will explore these representations in more detail later in this chapter.

- A. Using the notation of the example in the text, we write the two strings as follows:

```

0  0  2  7  C  8  F  8
00000000001001111100100011111000
*****
4  A  1  F  2      3      E      0
01001010000111110010001111100000
```

- B. With the second word shifted two positions to the right relative to the first, we find a sequence with 21 matching bits.
- C. We find all bits of the integer embedded in the floating-point number, except for the most significant bit having value 0. Such is the case for the example in the text as well. In addition, the floating-point number has some nonzero high-order bits that do not match those of the integer.

#### Solution to Problem 2.7 (page 85)

It prints 6D 6E 6F 70 71 72. Recall also that the library routine `strlen` does not count the terminating null character, and so `show_bytes` printed only through the character 'r'.

#### Solution to Problem 2.8 (page 87)

This problem is a drill to help you become more familiar with Boolean operations.

Operation	Result
$a$	[01001110]
$b$	[11100001]
$\sim a$	[10110001]
$\sim b$	[00011110]
$a \& b$	[01000000]
$a \mid b$	[11101111]
$a \wedge b$	[10101111]

**Solution to Problem 2.9 (page 89)**

This problem illustrates how Boolean algebra can be used to describe and reason about real-world systems. We can see that this color algebra is identical to the Boolean algebra over bit vectors of length 3.

- A. Colors are complemented by complementing the values of  $R$ ,  $G$ , and  $B$ . From this, we can see that white is the complement of black, yellow is the complement of blue, magenta is the complement of green, and cyan is the complement of red.
- B. We perform Boolean operations based on a bit-vector representation of the colors. From this we get the following:

$$\begin{array}{llll}
 \text{Blue (001)} & | & \text{Green (010)} & = \text{Cyan (011)} \\
 \text{Yellow (110)} & \& \text{Cyan (011)} & = \text{Green (010)} \\
 \text{Red (100)} & \sim & \text{Magenta (101)} & = \text{Blue (001)}
 \end{array}$$

**Solution to Problem 2.10 (page 90)**

This procedure relies on the fact that EXCLUSIVE-OR is commutative and associative, and that  $a \wedge a = 0$  for any  $a$ .

Step	*x	*y
Initially	$a$	$b$
Step 1	$a$	$a \wedge b$
Step 2	$a \wedge (a \wedge b) = (a \wedge a) \wedge b = b$	$a \wedge b$
Step 3	$b$	$b \wedge (a \wedge b) = (b \wedge b) \wedge a = a$

See Problem 2.11 for a case where this function will fail.

**Solution to Problem 2.11 (page 91)**

This problem illustrates a subtle and interesting feature of our inplace swap routine.

- A. Both `first` and `last` have value  $k$ , so we are attempting to swap the middle element with itself.
- B. In this case, arguments  $x$  and  $y$  to `inplace_swap` both point to the same location. When we compute  $*x \wedge *y$ , we get 0. We then store 0 as the middle element of the array, and the subsequent steps keep setting this element to 0. We can see that our reasoning in Problem 2.10 implicitly assumed that  $x$  and  $y$  denote different locations.
- C. Simply replace the test in line 4 of `reverse_array` to be `first < last`, since there is no need to swap the middle element with itself.

**Solution to Problem 2.12 (page 91)**

Here are the expressions:

- A.  $x \& 0xFF$
- B.  $x \wedge \sim 0xFF$
- C.  $x \mid 0xFF$

These expressions are typical of the kind commonly found in performing low-level bit operations. The expression  $\sim 0xFF$  creates a mask where the 8 least-significant bits equal 0 and the rest equal 1. Observe that such a mask will be generated regardless of the word size. By contrast, the expression  $0xFFFFF00$  would only work when data type `int` is 32 bits.

### Solution to Problem 2.13 (page 92)

These problems help you think about the relation between Boolean operations and typical ways that programmers apply masking operations. Here is the code:

```
/* Declarations of functions implementing operations bis and bic */
int bis(int x, int m);
int bic(int x, int m);

/* Compute x|y using only calls to functions bis and bic */
int bool_or(int x, int y) {
    int result = bis(x,y);
    return result;
}

/* Compute x^y using only calls to functions bis and bic */
int bool_xor(int x, int y) {
    int result = bis(bic(x,y), bic(y,x));
    return result;
}
```

The `bis` operation is equivalent to Boolean OR—a bit is set in `z` if either this bit is set in `x` or it is set in `m`. On the other hand, `bic(x, m)` is equivalent to  $x \& \sim m$ ; we want the result to equal 1 only when the corresponding bit of `x` is 1 and of `m` is 0.

Given that, we can implement `|` with a single call to `bis`. To implement  $\wedge$ , we take advantage of the property

$$x \wedge y = (x \& \sim y) \mid (\sim x \& y)$$

### Solution to Problem 2.14 (page 93)

This problem highlights the relation between bit-level Boolean operations and logical operations in C. A common programming error is to use a bit-level operation when a logical one is intended, or vice versa.

Expression	Value	Expression	Value
$a \& b$	0x44	$a \&\& b$	0x01
$a   b$	0x57	$a    b$	0x01
$\sim a   \sim b$	0xBB	$!a    !b$	0x00
$a \& !b$	0x00	$a \&\& \sim b$	0x01

**Solution to Problem 2.15 (page 93)**

The expression is  $!(x \hat{=} y)$ .

That is,  $x \hat{=} y$  will be zero if and only if every bit of  $x$  matches the corresponding bit of  $y$ . We then exploit the ability of  $!$  to determine whether a word contains any nonzero bit.

There is no real reason to use this expression rather than simply writing  $x == y$ , but it demonstrates some of the nuances of bit-level and logical operations.

**Solution to Problem 2.16 (page 94)**

This problem is a drill to help you understand the different shift operations.

x		a << 2		Logical a >> 3		Arithmetic a >> 3	
Hex	Binary	Binary	Hex	Binary	Hex	Binary	Hex
0xD4	[11010100]	[01010000]	0x50	[00011010]	0x1A	[11111010]	0xFA
0x64	[01100100]	[10010000]	0x90	[00001100]	0x0C	[11101100]	0xEC
0x72	[01110010]	[11001000]	0xC8	[00001110]	0x0E	[00001110]	0x0E
0x44	[01000100]	[00010000]	0x10	[00001000]	0x08	[11101000]	0xE8

**Solution to Problem 2.17 (page 101)**

In general, working through examples for very small word sizes is a very good way to understand computer arithmetic.

The unsigned values correspond to those in Figure 2.2. For the two's-complement values, hex digits 0 through 7 have a most significant bit of 0, yielding nonnegative values, while hex digits 8 through F have a most significant bit of 1, yielding a negative value.

Hexadecimal	Binary		
$\vec{x}$		$B2U_4(\vec{x})$	$B2T_4(\vec{x})$
0xA	[1010]	$2^3 + 2^1 = 10$	$-2^3 + 2^2 = -6$
0x1	[0001]	$2^0 = 1$	$2^0 = 1$
0xB	[1011]	$2^3 + 2^1 + 2^0 = 11$	$-2^3 + 2^1 + 2^0 = -5$
0x2	[0010]	$2^1 = 2$	$2^1 = 2$
0x7	[0111]	$2^2 + 2^1 + 2^0 = 7$	$2^2 + 2^1 + 2^0 = 7$
0xC	[1100]	$2^3 + 2^2 = 12$	$-2^3 + 2^2 = -4$

**Solution to Problem 2.18 (page 105)**

For a 32-bit word, any value consisting of 8 hexadecimal digits beginning with one of the digits 8 through f represents a negative number. It is quite common to see numbers beginning with a string of f's, since the leading bits of a negative number are all ones. You must look carefully, though. For example, the number 0x8048337 has only 7 digits. Filling this out with a leading zero gives 0x08048337, a positive number.

```

4004d0: 48 81 ec e0 02 00 00    sub    $0x2e0,%rsp           A. 736
4004d7: 48 8b 44 24 a8          mov    -0x58(%rsp),%rax      B. -88
4004dc: 48 03 47 28             add    0x28(%rdi),%rax       C. 40
4004e0: 48 89 44 24 d0          mov    %rax,-0x30(%rsp)      D. -48
4004e5: 48 8b 44 24 78          mov    0x78(%rsp),%rax       E. 120
4004ea: 48 89 87 88 00 00 00    mov    %rax,0x88(%rdi)      F. 136
4004f1: 48 8b 84 24 f8 01 00    mov    0x1f8(%rsp),%rax     G. 504
4004f8: 00
4004f9: 48 03 44 24 08          add    0x8(%rsp),%rax
4004fe: 48 89 84 24 c0 00 00    mov    %rax,0xc0(%rsp)      H. 192
400505: 00
400506: 48 8b 44 d4 b8          mov    -0x48(%rsp,%rdx,8),%rax I. -72

```

**Solution to Problem 2.19 (page 107)**

The functions  $T2U$  and  $U2T$  are very peculiar from a mathematical perspective. It is important to understand how they behave.

We solve this problem by reordering the rows in the solution of Problem 2.17 according to the two's-complement value and then listing the unsigned value as the result of the function application. We show the hexadecimal values to make this process more concrete.

$\vec{x}$ (hex)	$x$	$T2U_4(vecx)$
0xF	-1	15
0xB	-5	11
0xA	-6	10
0xC	-4	12
0x1	1	1
0x8	8	8

**Solution to Problem 2.20 (page 109)**

This exercise tests your understanding of Equation 2.5.

For the first four entries, the values of  $x$  are negative and  $T2U_4(x) = x + 2^4$ . For the remaining two entries, the values of  $x$  are nonnegative and  $T2U_4(x) = x$ .

**Solution to Problem 2.21 (page 112)**

This problem reinforces your understanding of the relation between two's-complement and unsigned representations, as well as the effects of the C promotion rules. Recall that  $TMin_{32}$  is  $-2,147,483,648$ , and that when cast to unsigned it

becomes 2,147,483,648. In addition, if either operand is unsigned, then the other operand will be cast to unsigned before comparing.

Expression	Type	Evaluation
<code>-2147483647-1 == 2147483648U</code>	Unsigned	1
<code>-2147483647-1 &lt; 2147483647</code>	Signed	1
<code>-2147483647-1U &lt; 2147483647</code>	Unsigned	0
<code>-2147483647-1 &lt; -2147483647</code>	Signed	1
<code>-2147483647-1U &lt; -2147483647</code>	Unsigned	1

#### Solution to Problem 2.22 (page 115)

This exercise provides a concrete demonstration of how sign extension preserves the numeric value of a two's-complement representation.

- A.  $[1100] \quad -2^3 + 2^2 = -8 + 4 = -4$   
 B.  $[11100] \quad -2^4 + 2^3 + 2^2 = -16 + 8 + 4 = -4$   
 C.  $[111100] \quad -2^5 + 2^4 + 2^3 + 2^2 = -32 + 16 + 8 + 4 = -4$

#### Solution to Problem 2.23 (page 116)

The expressions in these functions are common program “idioms” for extracting values from a word in which multiple bit fields have been packed. They exploit the zero-filling and sign-extending properties of the different shift operations. Note carefully the ordering of the cast and shift operations. In `fun1`, the shifts are performed on unsigned variable `w` and hence are logical. In `fun2`, shifts are performed after casting `w` to `int` and hence are arithmetic.

A.	w	fun1(w)	fun2(w)
	0x00000076	0x00000076	0x00000076
	0x87654321	0x00000021	0x00000021
	0x000000C9	0x000000C9	0xFFFFF8C9
	0xEDCBA987	0x00000087	0xFFFFF887

- B. Function `fun1` extracts a value from the low-order 8 bits of the argument, giving an integer ranging between 0 and 255. Function `fun2` extracts a value from the low-order 8 bits of the argument, but it also performs sign extension. The result will be a number between  $-128$  and  $127$ .

#### Solution to Problem 2.24 (page 118)

The effect of truncation is fairly intuitive for unsigned numbers, but not for two's-complement numbers. This exercise lets you explore its properties using very small word sizes.

Hex		Unsigned		Two's complement	
Original	Truncated	Original	Truncated	Original	Truncated
1	1	1	1	1	1
3	3	3	3	3	3
5	5	5	5	5	5
C	4	12	4	-4	4
E	6	14	6	-2	6

As Equation 2.9 states, the effect of this truncation on unsigned values is to simply find their residue, modulo 8. The effect of the truncation on signed values is a bit more complex. According to Equation 2.10, we first compute the modulo 8 residue of the argument. This will give values 0 through 7 for arguments 0 through 7, and also for arguments  $-8$  through  $-1$ . Then we apply function  $U2T_3$  to these residues, giving two repetitions of the sequences 0 through 3 and  $-4$  through  $-1$ .

#### Solution to Problem 2.25 (page 119)

This problem is designed to demonstrate how easily bugs can arise due to the implicit casting from signed to unsigned. It seems quite natural to pass parameter `length` as an unsigned, since one would never want to use a negative length. The stopping criterion `i <= length-1` also seems quite natural. But combining these two yields an unexpected outcome!

Since parameter `length` is unsigned, the computation `0 - 1` is performed using unsigned arithmetic, which is equivalent to modular addition. The result is then `UMax`. The `<=` comparison is also performed using an unsigned comparison, and since any number is less than or equal to `UMax`, the comparison always holds! Thus, the code attempts to access invalid elements of array `a`.

The code can be fixed either by declaring `length` to be an `int` or by changing the test of the for loop to be `i < length`.

#### Solution to Problem 2.26 (page 119)

This example demonstrates a subtle feature of unsigned arithmetic, and also the property that we sometimes perform unsigned arithmetic without realizing it. This can lead to very tricky bugs.

- For what cases will this function produce an incorrect result?* The function will incorrectly return 1 when `s` is shorter than `t`.
- Explain how this incorrect result comes about.* Since `strlen` is defined to yield an unsigned result, the difference and the comparison are both computed using unsigned arithmetic. When `s` is shorter than `t`, the difference `strlen(s) - strlen(t)` should be negative, but instead becomes a large, unsigned number, which is greater than 0.
- Show how to fix the code so that it will work reliably.* Replace the test with the following:

```
return strlen(s) > strlen(t);
```



**Solution to Problem 2.27 (page 125)**

This function is a direct implementation of the rules given to determine whether or not an unsigned addition overflows.

```
/* Determine whether arguments can be added without overflow */
int uadd_ok(unsigned x, unsigned y) {
    unsigned sum = x+y;
    return sum >= x;
}
```

**Solution to Problem 2.28 (page 125)**

This problem is a simple demonstration of arithmetic modulo 16. The easiest way to solve it is to convert the hex pattern into its unsigned decimal value. For nonzero values of  $x$ , we must have  $(-\frac{u}{4}x) + x = 16$ . Then we convert the complemented value back to hex.

$x$		$-\frac{u}{4}x$	
Hex	Decimal	Decimal	Hex
1	1	15	F
4	4	12	C
7	7	9	9
A	10	6	6
E	14	2	2

**Solution to Problem 2.29 (page 129)**

This problem is an exercise to make sure you understand two's-complement addition.

$x$	$y$	$x + y$	$x +_5^l y$	Case
-12	-15	-27	5	1
[10100]	[10001]	[100101]	[00101]	
-8	-8	-16	-16	2
[11000]	[11000]	[110000]	[10000]	
-9	8	-1	-1	2
[10111]	[01000]	[111111]	[11111]	
2	5	7	7	3
[00010]	[00101]	[000111]	[00111]	
12	4	16	-16	4
[01100]	[00100]	[010000]	[10000]	

**Solution to Problem 2.30 (page 130)**

This function is a direct implementation of the rules given to determine whether or not a two's-complement addition overflows.

```
/* Determine whether arguments can be added without overflow */
int tadd_ok(int x, int y) {
    int sum = x+y;
    int neg_over = x < 0 && y < 0 && sum >= 0;
    int pos_over = x >= 0 && y >= 0 && sum < 0;
    return !neg_over && !pos_over;
}
```

**Solution to Problem 2.31 (page 130)**

Your coworker could have learned, by studying Section 2.3.2, that two's-complement addition forms an abelian group, and so the expression  $(x+y)-x$  will evaluate to  $y$  regardless of whether or not the addition overflows, and that  $(x+y)-y$  will always evaluate to  $x$ .

**Solution to Problem 2.32 (page 130)**

This function will give correct values, except when  $y$  is  $TMin$ . In this case, we will have  $-y$  also equal to  $TMin$ , and so the call to function `tadd_ok` will indicate overflow when  $x$  is negative and no overflow when  $x$  is nonnegative. In fact, the opposite is true: `tsub_ok(x, TMin)` should yield 0 when  $x$  is negative and 1 when it is nonnegative.

One lesson to be learned from this exercise is that  $TMin$  should be included as one of the cases in any test procedure for a function.

**Solution to Problem 2.33 (page 131)**

This problem helps you understand two's-complement negation using a very small word size.

For  $w = 4$ , we have  $TMin_4 = -8$ . So  $-8$  is its own additive inverse, while other values are negated by integer negation.

$x$		$-_4^t x$	
Hex	Decimal	Decimal	Hex
2	2	-2	E
3	3	-3	D
9	-9	-9	7
B	-5	5	5
C	-4	4	4

The bit patterns are the same as for unsigned negation.

**Solution to Problem 2.34 (page 134)**

This problem is an exercise to make sure you understand two's-complement multiplication.

Mode	$x$		$y$		$x \cdot y$		Truncated $x \cdot y$	
Unsigned	4	[100]	5	[101]	20	[010100]	4	[100]
Two's complement	-4	[100]	-3	[101]	12	[001100]	-4	[100]
Unsigned	2	[010]	7	[111]	14	[001110]	6	[110]
Two's complement	2	[010]	-1	[111]	-2	[111110]	-2	[110]
Unsigned	6	[110]	6	[110]	36	[100100]	4	[100]
Two's complement	-2	[110]	-2	[110]	4	[000100]	-4	[100]

**Solution to Problem 2.35 (page 135)**

It is not realistic to test this function for all possible values of  $x$  and  $y$ . Even if you could run 10 billion tests per second, it would require over 58 years to test all combinations when data type `int` is 32 bits. On the other hand, it is feasible to test your code by writing the function with data type `short` or `char` and then testing it exhaustively.

Here's a more principled approach, following the proposed set of arguments:

1. We know that  $x \cdot y$  can be written as a  $2w$ -bit two's-complement number. Let  $u$  denote the unsigned number represented by the lower  $w$  bits, and  $v$  denote the two's-complement number represented by the upper  $w$  bits. Then, based on Equation 2.3, we can see that  $x \cdot y = v2^w + u$ .

We also know that  $u = T2U_w(p)$ , since they are unsigned and two's-complement numbers arising from the same bit pattern, and so by Equation 2.6, we can write  $u = p + p_{w-1}2^w$ , where  $p_{w-1}$  is the most significant bit of  $p$ . Letting  $t = v + p_{w-1}$ , we have  $x \cdot y = p + t2^w$ .

When  $t = 0$ , we have  $x \cdot y = p$ ; the multiplication does not overflow. When  $t \neq 0$ , we have  $x \cdot y \neq p$ ; the multiplication does overflow.

2. By definition of integer division, dividing  $p$  by nonzero  $x$  gives a quotient  $q$  and a remainder  $r$  such that  $p = x \cdot q + r$ , and  $|r| < |x|$ . (We use absolute values here, because the signs of  $x$  and  $r$  may differ. For example, dividing  $-7$  by  $2$  gives quotient  $-3$  and remainder  $-1$ .)
3. Suppose  $q = y$ . Then we have  $x \cdot y = x \cdot y + r + t2^w$ . From this, we can see that  $r + t2^w = 0$ . But  $|r| < |x| \leq 2^w$ , and so this identity can hold only if  $t = 0$ , in which case  $r = 0$ .

Suppose  $r = t = 0$ . Then we will have  $x \cdot y = x \cdot q$ , implying that  $y = q$ .

When  $x$  equals 0, multiplication does not overflow, and so we see that our code provides a reliable way to test whether or not two's-complement multiplication causes overflow.

**Solution to Problem 2.36 (page 135)**

With 64 bits, we can perform the multiplication without overflowing. We then test whether casting the product to 32 bits changes the value:

```

1  /* Determine whether the arguments can be multiplied
2     without overflow */
3  int tmult_ok(int x, int y) {
4     /* Compute product without overflow */
5     int64_t pll = (int64_t) x*y;
6     /* See if casting to int preserves value */
7     return pll == (int) pll;
8  }

```

Note that the casting on the right-hand side of line 5 is critical. If we instead wrote the line as

```
int64_t pll = x*y;
```

the product would be computed as a 32-bit value (possibly overflowing) and then sign extended to 64 bits.

#### Solution to Problem 2.37 (page 135)

- A. This change does not help at all. Even though the computation of `asize` will be accurate, the call to `malloc` will cause this value to be converted to a 32-bit unsigned number, and so the same overflow conditions will occur.
- B. With `malloc` having a 32-bit unsigned number as its argument, it cannot possibly allocate a block of more than  $2^{32}$  bytes, and so there is no point attempting to allocate or copy this much memory. Instead, the function should abort and return `NULL`, as illustrated by the following replacement to the original call to `malloc` (line 9):

```

uint64_t required_size = ele_cnt * (uint64_t) ele_size;
size_t request_size = (size_t) required_size;
if (required_size != request_size)
    /* Overflow must have occurred. Abort operation */
    return NULL;
void *result = malloc(request_size);
if (result == NULL)
    /* malloc failed */
    return NULL;

```

#### Solution to Problem 2.38 (page 138)

In Chapter 3, we will see many examples of the `LEA` instruction in action. The instruction is provided to support pointer arithmetic, but the C compiler often uses it as a way to perform multiplication by small constants.

For each value of  $k$ , we can compute two multiples:  $2^k$  (when  $b$  is 0) and  $2^k + 1$  (when  $b$  is  $a$ ). Thus, we can compute multiples 1, 2, 3, 4, 5, 8, and 9.

**Solution to Problem 2.39 (page 139)**

The expression simply becomes  $-(x \ll m)$ . To see this, let the word size be  $w$  so that  $n = w - 1$ . Form B states that we should compute  $(x \ll w) - (x \ll m)$ , but shifting  $x$  to the left by  $w$  will yield the value 0.

**Solution to Problem 2.40 (page 139)**

This problem requires you to try out the optimizations already described and also to supply a bit of your own ingenuity.

$K$	Shifts	Add/Subs	Expression
7	1	1	$(x \ll 3) - x$
30	4	3	$(x \ll 4) + (x \ll 3) + (x \ll 2) + (x \ll 1)$
28	2	1	$(x \ll 5) - (x \ll 2)$
55	2	2	$(x \ll 6) - (x \ll 3) - x$

Observe that the fourth case uses a modified version of form B. We can view the bit pattern [11011] as having a run of 6 ones with a zero in the middle, and so we apply the rule for form B, but then we subtract the term corresponding to the middle zero bit.

**Solution to Problem 2.41 (page 139)**

Assuming that addition and subtraction have the same performance, the rule is to choose form A when  $n = m$ , either form when  $n = m + 1$ , and form B when  $n > m + 1$ .

The justification for this rule is as follows. Assume first that  $m > 0$ . When  $n = m$ , form A requires only a single shift, while form B requires two shifts and a subtraction. When  $n = m + 1$ , both forms require two shifts and either an addition or a subtraction. When  $n > m + 1$ , form B requires only two shifts and one subtraction, while form A requires  $n - m + 1 > 2$  shifts and  $n - m > 1$  additions. For the case of  $m = 0$ , we get one fewer shift for both forms A and B, and so the same rules apply for choosing between the two.

**Solution to Problem 2.42 (page 143)**

The only challenge here is to compute the bias without any testing or conditional operations. We use the trick that the expression  $x \gg 31$  generates a word with all ones if  $x$  is negative, and all zeros otherwise. By masking off the appropriate bits, we get the desired bias value.

```
int div16(int x) {
    /* Compute bias to be either 0 (x >= 0) or 15 (x < 0) */
    int bias = (x >> 31) & 0xF;
    return (x + bias) >> 4;
}
```

**Solution to Problem 2.43 (page 143)**

We have found that people have difficulty with this exercise when working directly with assembly code. It becomes more clear when put in the form shown in `optarith`.

We can see that  $M$  is 31;  $x * M$  is computed as  $(x \ll 5) - x$ .

We can see that  $N$  is 8; a bias value of 7 is added when  $y$  is negative, and the right shift is by 3.

**Solution to Problem 2.44 (page 144)**

These “C puzzle” problems provide a clear demonstration that programmers must understand the properties of computer arithmetic:

- A.  $(x > 0) \mid\mid (x - 1 < 0)$   
*False.* Let  $x$  be  $-2,147,483,648$  ( $TMin_{32}$ ). We will then have  $x - 1$  equal to  $2,147,483,647$  ( $TMax_{32}$ ).
- B.  $(x \& 7) != 7 \mid\mid (x \ll 29 < 0)$   
*True.* If  $(x \& 7) != 7$  evaluates to 0, then we must have bit  $x_2$  equal to 1. When shifted left by 29, this will become the sign bit.
- C.  $(x * x) \geq 0$   
*False.* When  $x$  is  $65,535$  ( $0xFFFF$ ),  $x * x$  is  $-131,071$  ( $0xFFFFE0001$ ).
- D.  $x < 0 \mid\mid -x \leq 0$   
*True.* If  $x$  is nonnegative, then  $-x$  is nonpositive.
- E.  $x > 0 \mid\mid -x \geq 0$   
*False.* Let  $x$  be  $-2,147,483,648$  ( $TMin_{32}$ ). Then both  $x$  and  $-x$  are negative.
- F.  $x + y == uy + ux$   
*True.* Two’s-complement and unsigned addition have the same bit-level behavior, and they are commutative.
- G.  $x * \sim y + uy * ux == -x$   
*True.*  $\sim y$  equals  $-y - 1$ .  $uy * ux$  equals  $x * y$ . Thus, the left-hand side is equivalent to  $x * -y - x + x * y$ .

**Solution to Problem 2.45 (page 147)**

Understanding fractional binary representations is an important step to understanding floating-point encodings. This exercise lets you try out some simple examples.

$\frac{1}{8}$	0.001	0.125
$\frac{3}{4}$	0.11	0.75
$\frac{25}{16}$	1.1001	1.5625
$\frac{43}{16}$	10.1011	2.6875
$\frac{9}{8}$	1.001	1.125
$\frac{47}{8}$	101.111	5.875
$\frac{51}{16}$	11.0011	3.1875



### Solution to Problem 2.48 (page 155)

Hexadecimal 0x359141 is equivalent to binary [1101011001000101000001]. Shifting this right 21 places gives  $1.101011001000101000001_2 \times 2^{21}$ . We form the fraction field by dropping the leading 1 and adding two zeros, giving

$[10101100100010100000100]$

The exponent is formed by adding bias 127 to 21, giving 148 (binary [10010100]). We combine this with a sign field of 0 to give a binary representation

[01001010010101100100010100000100]

We see that the matching bits in the two representations correspond to the low-order bits of the integer, up to the most significant bit equal to 1 matching the high-order 21 bits of the fraction:

```

      0  0  3  5  9  1  4  1
00000000001101011001000101000001
      *****
      4  A  5  6  4  5  0  4
01001010010101100100010100000100

```

### Solution to Problem 2.49 (page 156)

This exercise helps you think about what numbers cannot be represented exactly in floating point.

- A. The number has binary representation 1, followed by  $n$  zeros, followed by 1, giving value  $2^{n+1} + 1$ .
- B. When  $n = 23$ , the value is  $2^{24} + 1 = 16,777,217$ .

### Solution to Problem 2.50 (page 157)

Performing rounding by hand helps reinforce the idea of round-to-even with binary numbers.

Original		Rounded	
$10.111_2$	$2\frac{7}{8}$	11.0	3
$11.010_2$	$3\frac{1}{4}$	11.0	3
$11.000_2$	3	11.0	3
$10.110_2$	$2\frac{3}{4}$	11.0	3

### Solution to Problem 2.51 (page 158)

- A. Looking at the nonterminating sequence for  $\frac{1}{10}$ , we see that the 2 bits to the right of the rounding position are 1, so a better approximation to  $\frac{1}{10}$  would be obtained by incrementing  $x$  to get  $x' = 0.00011001100110011001101_2$ , which is larger than 0.1.
- B. We can see that  $x' - 0.1$  has binary representation

0.00000000000000000000000000000000[1100]



Comparing this to the binary representation of  $\frac{1}{10}$ , we can see that it is  $2^{-22} \times \frac{1}{10}$ , which is around  $2.38 \times 10^{-8}$ .

- C.  $2.38 \times 10^{-8} \times 100 \times 60 \times 60 \times 10 \approx 0.086$  seconds, a factor of 4 less than the error in the Patriot system.
- D.  $0.086 \times 2,000 \approx 171$  meters.

#### Solution to Problem 2.52 (page 158)

This problem tests a lot of concepts about floating-point representations, including the encoding of normalized and denormalized values, as well as rounding.

Format A		Format B		Comments
Bits	Value	Bits	Value	
011 0000	1	0111 000	1	
101 1110	$\frac{15}{2}$	1001 111	$\frac{15}{2}$	
010 1001	$\frac{25}{32}$	0110 100	$\frac{3}{4}$	Round down
110 1111	$\frac{31}{2}$	1011 000	16	Round up
000 0001	$\frac{1}{64}$	0001 000	$\frac{1}{64}$	Denorm $\rightarrow$ norm

#### Solution to Problem 2.53 (page 161)

In general, it is better to use a library macro rather than inventing your own code. This code seems to work on a variety of machines, however.

We assume that the value 1e400 overflows to infinity.

```
#define POS_INFINITY 1e400
#define NEG_INFINITY (-POS_INFINITY)
#define NEG_ZERO (-1.0/POS_INFINITY)
```

#### Solution to Problem 2.54 (page 161)

Exercises such as this one help you develop your ability to reason about floating-point operations from a programmer's perspective. Make sure you understand each of the answers.

- A. `x == (int)(double) x`  
Yes, since double has greater precision and range than int.
- B. `x == (int)(float) x`  
No. For example, when x is *TMax*.
- C. `d == (double)(float) d`  
No. For example, when d is 1e40, we will get  $+\infty$  on the right.
- D. `f == (float)(double) f`  
Yes, since double has greater precision and range than float.
- E. `f == -(-f)`  
Yes, since a floating-point number is negated by simply inverting its sign bit.

F.  $1.0/2 == 1/2.0$

Yes, the numerators and denominators will both be converted to floating-point representations before the division is performed.

G.  $d*d \geq 0.0$

Yes, although it may overflow to  $+\infty$ .

H.  $(f+d)-f == d$

No. For example, when  $f$  is  $1.0e20$  and  $d$  is  $1.0$ , the expression  $f+d$  will be rounded to  $1.0e20$ , and so the expression on the left-hand side will evaluate to  $0.0$ , while the right-hand side will be  $1.0$ .